

Neural Network Architecture

Borrowing the idea from ResNet[1, 2], I use two mechanisms here, the skip connection and the pre-activation, to construct my own neural network (See Figure 1). It is supposed that we are training a multi-class classifier with the cross-entropy loss. It can be seen from Figure 1 that all the intermediate layers, except activation functions, are fully connected layers.




Figure 1: The architecture of this self-designed neural network.

A forward pass of a single input sample \mathbf{x} is as follows:

$$\begin{aligned}
 \mathbf{x} &\in \mathbb{R}^D \\
 \mathbf{z} &= \text{ReLU}(\mathbf{x}) \in \mathbb{R}^D \\
 \mathbf{h} &= W_1 \mathbf{z} \in \mathbb{R}^D \\
 \mathbf{s} &= \mathbf{h} + \mathbf{x} \in \mathbb{R}^D \\
 \mathbf{t} &= W_2 \mathbf{s} \in \mathbb{R}^M \\
 \mathbf{p} &= \text{ReLU}(\mathbf{t}) \in \mathbb{R}^M \\
 \mathbf{q} &= W_3 \mathbf{p} \in \mathbb{R}^M \\
 \mathbf{r} &= \mathbf{q} + \mathbf{t} \in \mathbb{R}^M \\
 \boldsymbol{\theta} &= W_4 \mathbf{r} \in \mathbb{R}^C \\
 \hat{\mathbf{y}} &= \text{softmax}(\boldsymbol{\theta}) \in \mathbb{R}^C \\
 J &= \text{CrossEntropy}(\mathbf{y}, \hat{\mathbf{y}}) \in \mathbb{R}
 \end{aligned}$$

where $W_1 \in \mathbb{R}^{D \times D}$, $W_2 \in \mathbb{R}^{M \times D}$, $W_3 \in \mathbb{R}^{M \times M}$, $W_4 \in \mathbb{R}^{C \times M}$ are weight matrices. $\mathbf{y} \in \mathbb{R}^C$ is the true label vector associated with \mathbf{x} , that is, a one-hot vector indicating which class \mathbf{x} belongs to. The softmax activation of the j -th output unit is

$$\hat{y}_j = \text{softmax}(\theta_j) = \frac{e^{\theta_j}}{\sum_{j=1}^C e^{\theta_j}}.$$

The cross-entropy loss for this single input sample \mathbf{x} is

$$J = -\sum_{j=1}^C y_j \log \hat{y}_j = -\mathbf{y}^T \log \hat{\mathbf{y}}.$$

Gradients Derivation

First, let us start with calculating the common parts shared by all four target gradients. we can see that

$$\frac{\partial J}{\partial \hat{\mathbf{y}}} = \frac{\partial (-\mathbf{y}^T \log \hat{\mathbf{y}})}{\partial \hat{\mathbf{y}}} = -\left[\frac{y_1}{\hat{y}_1}, \dots, \frac{y_C}{\hat{y}_C}\right] = -\left(\frac{\mathbf{y}}{\hat{\mathbf{y}}}\right)^T \in \mathbb{R}^{1 \times C}$$

where the division in $\frac{\mathbf{y}}{\hat{\mathbf{y}}}$ is an element-wise operator.

For the gradient of softmax function, we look at a single entry first:

$$\frac{\partial \hat{y}_i}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \frac{e^{\theta_i}}{\sum_{k=1}^C e^{\theta_k}} = \begin{cases} \frac{e^{\theta_j} (\sum_{k=1}^C e^{\theta_k} - e^{\theta_j})}{(\sum_{k=1}^C e^{\theta_k})^2} = \hat{y}_j (1 - \hat{y}_j), i = j \\ -\frac{e^{\theta_i} e^{\theta_j}}{(\sum_{k=1}^C e^{\theta_k})^2} = -\hat{y}_i \hat{y}_j, i \neq j \end{cases}.$$

By chain rule, we can compute

$$\begin{aligned} \frac{\partial \mathcal{J}}{\partial \theta_j} &= \sum_{i=1}^C \frac{\partial \mathcal{J}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial \theta_j} = - \sum_{i=1}^C \frac{y_i}{\hat{y}_i} \frac{\partial \hat{y}_i}{\partial \theta_j} = - \left(\frac{y_j}{\hat{y}_j} \frac{\partial \hat{y}_j}{\partial \theta_j} + \sum_{i \neq j} \frac{y_i}{\hat{y}_i} \frac{\partial \hat{y}_i}{\partial \theta_j} \right) \\ &= -\frac{y_j}{\hat{y}_j} \hat{y}_j (1 - \hat{y}_j) + \sum_{i \neq j} \frac{y_i}{\hat{y}_i} \hat{y}_i \hat{y}_j = -y_j (1 - \hat{y}_j) + \hat{y}_j \sum_{i \neq j} y_i \\ &= -y_j + \hat{y}_j \sum_{i=1}^C y_i. \end{aligned}$$

Recall that \mathbf{y} is a one-hot vector, that is, $\sum_{i=1}^C y_i = 1$, we can get

$$\frac{\partial \mathcal{J}}{\partial \theta_j} = \hat{y}_j - y_j$$

and

$$\frac{\partial \mathcal{J}}{\partial \boldsymbol{\theta}} = (\hat{\mathbf{y}} - \mathbf{y})^T \in \mathbb{R}^{1 \times C}.$$

Now we examine the gradient of ReLU function. Recall that $\text{ReLU}(x) = \max(x, 0)$, by extending that its derivative at $x = 0$ equals 0, we get

$$\text{ReLU}'(x) = \begin{cases} 1, x > 0 \\ 0, \text{otherwise} \end{cases} = \mathbb{I}_{x>0}(x) \in \{0,1\}$$

where $\mathbb{I}_{x>0}(x)$ is an indicator function. Furthermore, we can extend ReLU function to a vector-valued function by applying ReLU to each entry of an input vector $\mathbf{x} \in \mathbb{R}^D$ and calculate its gradient as

$$\frac{\partial \text{ReLU}(\mathbf{x})}{\partial \mathbf{x}} = \text{diag}(\mathbb{I}_{x_1>0}(x_1), \dots, \mathbb{I}_{x_D>0}(x_D)) \in \mathbb{R}^{D \times D}.$$

It is easy to see that W_1, W_2, W_3 and W_4 contain all the network parameters. So we just need to derive their gradients to update the whole neural network.

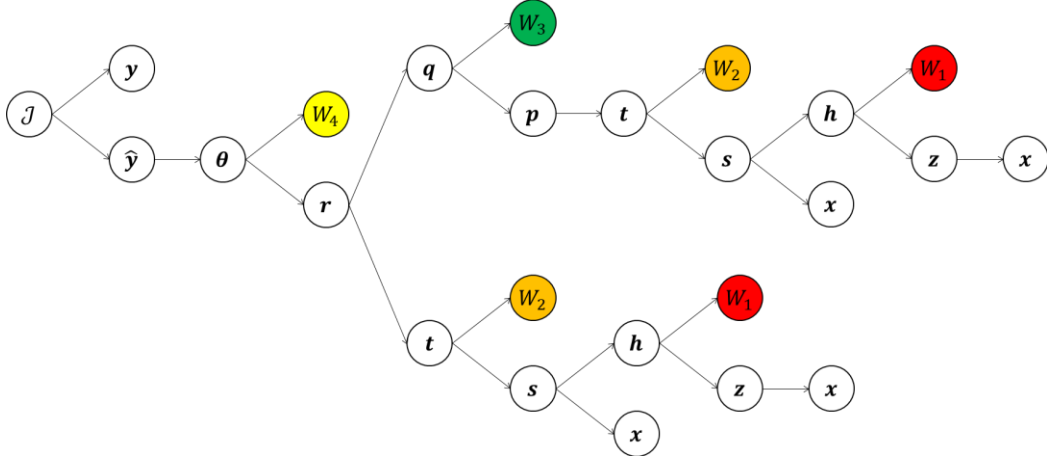


Figure 2: The computation graph of this self-designed neural network.

However, due to the complexity of this neural network, it is hard to capture directly how the gradients flow just by observing its architecture figure. So I show the computation graph of it to make the process of gradients derivation more clearly (See Figure 2).

It can be seen obviously from Figure 2 that there is only one path to approach W_3 and W_4 , but two to approach W_1 and W_2 . With the help of this computation graph, we can write down the gradients for all the network parameters by using the chain rule as

$$\begin{aligned}\frac{\partial J}{\partial W_4} &= \frac{\partial J}{\partial \theta} \frac{\partial \theta}{\partial W_4} \\ \frac{\partial J}{\partial W_3} &= \frac{\partial J}{\partial \mathbf{r}} \frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial W_3} = \frac{\partial J}{\partial \theta} \frac{\partial \theta}{\partial \mathbf{r}} \frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial W_3} \\ \frac{\partial J}{\partial W_2} &= \frac{\partial J}{\partial \mathbf{r}} \left(\left(\frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial \mathbf{t}} \right)_1 \frac{\partial \mathbf{t}}{\partial W_2} + \left(\frac{\partial \mathbf{r}}{\partial \mathbf{t}} \right)_2 \frac{\partial \mathbf{t}}{\partial W_2} \right) \\ \frac{\partial J}{\partial W_1} &= \frac{\partial J}{\partial \mathbf{r}} \left(\left(\frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial \mathbf{t}} \right)_1 \frac{\partial \mathbf{t}}{\partial \mathbf{s}} \frac{\partial \mathbf{s}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial W_1} + \left(\frac{\partial \mathbf{r}}{\partial \mathbf{t}} \right)_2 \frac{\partial \mathbf{t}}{\partial \mathbf{s}} \frac{\partial \mathbf{s}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial W_1} \right).\end{aligned}$$

Subscript 1 means the gradient flows through the upper branch of the computation graph and subscript 2 the lower. We use subscripts here to avoid confusion, that is, $\left(\frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial \mathbf{t}} \right)_1 \neq \left(\frac{\partial \mathbf{r}}{\partial \mathbf{t}} \right)_2$.

It is noticed that $\frac{\partial J}{\partial \mathbf{r}}$, $\left(\frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial \mathbf{t}} \right)_1$, $\frac{\partial \mathbf{t}}{\partial \mathbf{s}} \frac{\partial \mathbf{s}}{\partial \mathbf{h}}$ and $\left(\frac{\partial \mathbf{r}}{\partial \mathbf{t}} \right)_2$ are used multiple times, so it is convenient to calculate them firstly:

$$\begin{aligned}\frac{\partial J}{\partial \mathbf{r}} &= \frac{\partial J}{\partial \theta} \frac{\partial \theta}{\partial \mathbf{r}} = (\hat{\mathbf{y}} - \mathbf{y})^T W_4 \in \mathbb{R}^{1 \times M} \\ \left(\frac{\partial \mathbf{r}}{\partial \mathbf{q}} \frac{\partial \mathbf{q}}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial \mathbf{t}} \right)_1 &= I_M W_3 \text{diag}(\mathbb{I}_{t_1 > 0}(t_1), \dots, \mathbb{I}_{t_M > 0}(t_M)) = W_3 \Lambda_{\mathbf{t}} \in \mathbb{R}^{M \times M} \\ \frac{\partial \mathbf{t}}{\partial \mathbf{s}} \frac{\partial \mathbf{s}}{\partial \mathbf{h}} &= W_2 I_D = W_2 \in \mathbb{R}^{M \times D} \\ \left(\frac{\partial \mathbf{r}}{\partial \mathbf{t}} \right)_2 &= I_M \in \mathbb{R}^{M \times M}\end{aligned}$$

where $\Lambda_{\mathbf{t}} = \text{diag}(\mathbb{I}_{t_1 > 0}(t_1), \dots, \mathbb{I}_{t_M > 0}(t_M))$.

Now we can rewrite the target gradients as

$$\begin{aligned}\frac{\partial J}{\partial W_4} &= (\hat{\mathbf{y}} - \mathbf{y})^T \frac{\partial \theta}{\partial W_4} \\ \frac{\partial J}{\partial W_3} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 \frac{\partial \mathbf{q}}{\partial W_3} \\ \frac{\partial J}{\partial W_2} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_3 \Lambda_{\mathbf{t}} \frac{\partial \mathbf{t}}{\partial W_2} + (\hat{\mathbf{y}} - \mathbf{y})^T W_4 \frac{\partial \mathbf{t}}{\partial W_2} \\ \frac{\partial J}{\partial W_1} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_3 \Lambda_{\mathbf{t}} W_2 \frac{\partial \mathbf{h}}{\partial W_1} + (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_2 \frac{\partial \mathbf{h}}{\partial W_1}.\end{aligned}$$

Then we continue to calculate the rest parts. For the gradient $\frac{\partial J}{\partial W_4}$, let us start by analyzing

the gradient $\frac{\partial \theta}{\partial W_4}$. By definition, this gradient is the collection of the partial derivatives:

$$\frac{\partial \boldsymbol{\theta}}{\partial W_4} = \begin{bmatrix} \frac{\partial \theta_1}{\partial W_4} \\ \vdots \\ \frac{\partial \theta_c}{\partial W_4} \end{bmatrix} \in \mathbb{R}^{C \times C \times M}, \quad \frac{\partial \theta_i}{\partial W_4} \in \mathbb{R}^{C \times M}.$$

We can write W_4 in this form:

$$W_4 = \begin{bmatrix} \mathbf{w}_1^{(4)} \\ \vdots \\ \mathbf{w}_c^{(4)} \end{bmatrix}$$

where $\mathbf{w}_i^{(4)} \in \mathbb{R}^{1 \times M}$ is the i -th row of W_4 . So we can explicitly write out $\frac{\partial \theta_i}{\partial W_4}$ as

$$\frac{\partial \theta_i}{\partial W_4} = \frac{\partial}{\partial W_4} \mathbf{w}_i^{(4)} \mathbf{r} = \begin{bmatrix} \mathbf{0}^T \\ \vdots \\ \mathbf{0}^T \\ \mathbf{r}^T \\ \mathbf{0}^T \\ \vdots \\ \mathbf{0}^T \end{bmatrix} \in \mathbb{R}^{C \times M}$$

that is, a matrix with i -th row being \mathbf{r}^T and other rows being $\mathbf{0}^T$. Combining them together, we can get

$$\frac{\partial \mathcal{J}}{\partial W_4} = (\hat{\mathbf{y}} - \mathbf{y})^T \frac{\partial \boldsymbol{\theta}}{\partial W_4} = (\hat{\mathbf{y}} - \mathbf{y})^T \begin{bmatrix} \frac{\partial \theta_1}{\partial W_4} \\ \vdots \\ \frac{\partial \theta_c}{\partial W_4} \end{bmatrix} = \begin{bmatrix} (\hat{\mathbf{y}} - \mathbf{y})^T \frac{\partial \theta_1}{\partial W_4} \\ \vdots \\ (\hat{\mathbf{y}} - \mathbf{y})^T \frac{\partial \theta_c}{\partial W_4} \end{bmatrix} = \begin{bmatrix} (\hat{y}_1 - y_1) \mathbf{r}^T \\ \vdots \\ (\hat{y}_c - y_c) \mathbf{r}^T \end{bmatrix} = (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{r}^T.$$

Similarly, we can get that

$$\begin{aligned} \frac{\partial \mathcal{J}}{\partial W_3} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 \frac{\partial \mathbf{q}}{\partial W_3} = W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{p}^T \\ \frac{\partial \mathcal{J}}{\partial W_2} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_3 \Lambda_t \frac{\partial \mathbf{t}}{\partial W_2} + (\hat{\mathbf{y}} - \mathbf{y})^T W_4 \frac{\partial \mathbf{t}}{\partial W_2} \\ &= \Lambda_t W_3^T W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{s}^T + W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{s}^T \\ \frac{\partial \mathcal{J}}{\partial W_1} &= (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_3 \Lambda_t W_2 \frac{\partial \mathbf{h}}{\partial W_1} + (\hat{\mathbf{y}} - \mathbf{y})^T W_4 W_2 \frac{\partial \mathbf{h}}{\partial W_1} \\ &= W_2^T \Lambda_t W_3^T W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{z}^T + W_2^T W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{z}^T. \end{aligned}$$

At the end, we show the final results as follows:

$$\begin{aligned} \frac{\partial \mathcal{J}}{\partial W_4} &= (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{r}^T \\ \frac{\partial \mathcal{J}}{\partial W_3} &= W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{p}^T \\ \frac{\partial \mathcal{J}}{\partial W_2} &= (\Lambda_t W_3^T + I_M) W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{s}^T \\ \frac{\partial \mathcal{J}}{\partial W_1} &= W_2^T (\Lambda_t W_3^T + I_M) W_4^T (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{z}^T. \end{aligned}$$

Training Equations

In real practice, we never use just a single input to update the neural network. Instead, we train a neural network with a training set in most cases. To make this assignment more practical, we suppose that we use the training set $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$ to train the neural network.

As a result, the final loss function needs to be modified as

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \mathcal{J}_i(\mathbf{x}_i, \mathbf{y}_i)$$

where $\mathcal{J}_i(\mathbf{x}_i, \mathbf{y}_i)$ is the cross-entropy loss w.r.t a single input $(\mathbf{x}_i, \mathbf{y}_i)$, whose gradients are shown in the above section. And the gradients of it are

$$\frac{\partial \mathcal{L}}{\partial W_4} = \frac{1}{N} \sum_{i=1}^N \frac{\partial \mathcal{J}_i}{\partial W_4} = \frac{1}{N} \sum_{i=1}^N (\hat{\mathbf{y}}_i - \mathbf{y}_i) \mathbf{r}_i^T$$

$$\frac{\partial \mathcal{L}}{\partial W_3} = \frac{1}{N} \sum_{i=1}^N \frac{\partial \mathcal{J}_i}{\partial W_3} = \frac{1}{N} W_4^T \sum_{i=1}^N (\hat{\mathbf{y}}_i - \mathbf{y}_i) \mathbf{p}_i^T$$

$$\frac{\partial \mathcal{L}}{\partial W_2} = \frac{1}{N} \sum_{i=1}^N \frac{\partial \mathcal{J}_i}{\partial W_2} = \frac{1}{N} \sum_{i=1}^N (\Lambda_{t_i} W_3^T + I_M) W_4^T (\hat{\mathbf{y}}_i - \mathbf{y}_i) \mathbf{s}_i^T$$

$$\frac{\partial \mathcal{L}}{\partial W_1} = \frac{1}{N} \sum_{i=1}^N \frac{\partial \mathcal{J}_i}{\partial W_1} = \frac{1}{N} W_2^T \sum_{i=1}^N (\Lambda_{t_i} W_3^T + I_M) W_4^T (\hat{\mathbf{y}}_i - \mathbf{y}_i) \mathbf{z}_i^T.$$

For the training process, we start by randomly generating $W_1 \sim W_4$. At step t , we use the following training equations to update the parameters:

$$W_4^t \leftarrow W_4^{t-1} - \alpha \frac{\partial \mathcal{L}}{\partial W_4}$$

$$W_3^t \leftarrow W_3^{t-1} - \alpha \frac{\partial \mathcal{L}}{\partial W_3}$$

$$W_2^t \leftarrow W_2^{t-1} - \alpha \frac{\partial \mathcal{L}}{\partial W_2}$$

$$W_1^t \leftarrow W_1^{t-1} - \alpha \frac{\partial \mathcal{L}}{\partial W_1}$$

where α is the learning rate.

Reference

[1] He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition. 2016.

[2] He, Kaiming, et al. "Identity mappings in deep residual networks." European conference on computer vision. Springer, Cham, 2016.